

# Module 2: NCR Ride Bookings - Complete Project Guide

## Data Manipulation & Exploration with Pandas and NumPy

DS Club | Hands-On Project

**Duration:** 4 Sessions (90 minutes each)

**Dataset:** 150,000 ride booking records from NCR region

---

### Table of Contents

- [Project Overview](#)
  - [Session 1: Data Loading & Initial Exploration](#)
  - [Session 2: Data Cleaning & Preparation](#)
  - [Session 3: Exploratory Data Analysis \(EDA\)](#)
  - [Session 4: Advanced Analysis & Insights](#)
  - [Complete Code Reference](#)
- 

## Project Overview

### Business Context

You are a **Data Analyst** at a ride-hailing company operating in the NCR (National Capital Region). The company has been experiencing:

- Increasing ride cancellations
- Fluctuating revenues across different locations
- Unclear patterns in customer and driver behavior

**Your Mission:** Analyze 150,000 ride booking records to uncover insights that will help the business:

1. Reduce cancellation rates
2. Identify profitable locations
3. Understand peak demand times
4. Improve customer and driver satisfaction

### Dataset Information

**File:** ncr\_ride\_bookings.csv

**Size:** 150,000 rows × 21 columns

**Period:** January - March 2022

**Regions:** Delhi, Gurgaon, Noida, and surrounding NCR areas

#### Columns:

1. **Date** - Date of booking
2. **Time** - Time of booking
3. **Booking ID** - Unique identifier for each booking
4. **Booking Status** - Status of the ride (Success, Cancelled, etc.)
5. **Customer ID** - Unique customer identifier

6. **Vehicle Type** - Type of vehicle (Auto, Bike, Prime Sedan, Prime SUV)
  7. **Pickup Location** - Starting point of ride
  8. **Drop Location** - Destination
  9. **Avg VTAT** - Average Vehicle Turn Around Time
  10. **Avg CTAT** - Average Customer Turn Around Time
  11. **Cancelled Rides by Customer** - Number of cancellations by customer
  12. **Reason for cancelling by Customer** - Why customer cancelled
  13. **Cancelled Rides by Driver** - Number of cancellations by driver
  14. **Driver Cancellation Reason** - Why driver cancelled
  15. **Incomplete Rides** - Number of incomplete rides
  16. **Incomplete Rides Reason** - Why ride was incomplete
  17. **Booking Value** - Fare amount in local currency
  18. **Ride Distance** - Distance in kilometers
  19. **Driver Ratings** - Rating given to driver (1-5)
  20. **Customer Rating** - Rating given to customer (1-5)
  21. **Payment Method** - How customer paid (Cash, Card, Wallet, UPI)
- 

## Session 1: Data Loading & Initial Exploration

**Duration:** 90 minutes

**Goal:** Load the data and understand its structure

---

### Part 1: Setting Up Your Workspace (10 minutes)

#### Step 1.1: Import Required Libraries



python

```
# Import libraries for data manipulation
import pandas as pd      # For working with tabular data
import numpy as np        # For numerical operations
import matplotlib.pyplot as plt # For creating visualizations
import seaborn as sns     # For statistical visualizations

# Display settings
pd.set_option('display.max_columns', None) # Show all columns
pd.set_option('display.width', None)       # Don't truncate display
pd.set_option('display.max_rows', 100)      # Show up to 100 rows

print("✅ Libraries imported successfully!")
```

**What Each Library Does:**

- **pandas (pd):** Think of it as Excel on steroids - lets you work with tables of data (called DataFrames)
- **numpy (np):** For math operations, especially on arrays of numbers
- **matplotlib (plt):** Creates charts and graphs
- **seaborn (sns):** Makes matplotlib prettier and easier to use

The `pd.set_option()` commands:

- Help us see more data when we print DataFrames
- Makes debugging easier
- You can adjust these based on your screen size

---

## Question 2: Revenue by Vehicle Type



python

```

# Analyze revenue patterns by vehicle type
vehicle_revenue = df_clean.groupby('Vehicle Type').agg({
    'Booking Value': ['sum', 'mean', 'median', 'count'],
    'Ride Distance': 'mean'
}).round(2)

# Flatten column names
vehicle_revenue.columns = ['Total_Revenue', 'Avg_Booking', 'Median_Booking', 'Num_Rides', 'Avg_Distance']
vehicle_revenue = vehicle_revenue.sort_values('Total_Revenue', ascending=False)

print("Revenue Analysis by Vehicle Type:")
print(vehicle_revenue)

# Visualize - Multiple subplots
fig, axes = plt.subplots(1, 2, figsize=(15, 6))

# Plot 1: Total Revenue
vehicle_revenue['Total_Revenue'].plot(kind='bar', ax=axes[0], color='teal')
axes[0].set_title('Total Revenue by Vehicle Type', fontsize=14, fontweight='bold')
axes[0].set_ylabel('Revenue (₹)')
axes[0].set_xlabel('Vehicle Type')
axes[0].tick_params(axis='x', rotation=45)
axes[0].grid(True, alpha=0.3, axis='y')

# Plot 2: Average Booking Value
vehicle_revenue['Avg_Booking'].plot(kind='bar', ax=axes[1], color='coral')
axes[1].set_title('Average Booking Value by Vehicle Type', fontsize=14, fontweight='bold')
axes[1].set_ylabel('Avg Booking (₹)')
axes[1].set_xlabel('Vehicle Type')
axes[1].tick_params(axis='x', rotation=45)
axes[1].grid(True, alpha=0.3, axis='y')

plt.tight_layout()
plt.show()

```

\*\*Business Step 1.2: Load the Dataset



python

```
# Load the CSV file into a DataFrame
df = pd.read_csv('data/ncr_ride_bookings.csv')

# Display success message
print("✅ Dataset loaded successfully!")
print(f"Shape: {df.shape}")
```

### Understanding `pd.read_csv()`:

- **What it does:** Reads a CSV (Comma Separated Values) file and converts it to a DataFrame
- **Parameters:**
  - 'data/ncr\_ride\_bookings.csv' - Path to your file
  - You can add: `encoding='utf-8'` if you get encoding errors
  - You can add: `parse_dates=['Date']` to automatically convert dates

### Understanding `df.shape`:

- Returns a tuple: (`number_of_rows, number_of_columns`)
- Example output: (150000, 21) means 150,000 rows and 21 columns

---

## Part 2: First Look at the Data (20 minutes)

### Step 2.1: View Sample Data



python

```
# Display first 5 rows
print("First 5 rows of the dataset:")
df.head()
```

### Understanding `df.head()`:

- **What it does:** Shows the first 5 rows of your DataFrame
- **Why use it:** Quickly see what your data looks like
- **Variations:**
  - `df.head(10)` - Show first 10 rows
  - `df.tail()` - Show last 5 rows
  - `df.tail(20)` - Show last 20 rows
- **Returns:** A DataFrame with n rows



python

```
# Display last 5 rows
print("Last 5 rows of the dataset:")
df.tail()
```

## Why look at both head and tail?

- Check if data is consistent throughout
- See if there are any patterns in how data is organized
- Detect if data collection changed over time

---

## Step 2.2: Understanding Data Structure



python

```
# Get detailed information about the DataFrame
df.info()
```

### Understanding df.info() Output:



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150000 entries, 0 to 149999
Data columns (total 21 columns):
 #  Column           Non-Null Count  Dtype  
 --- 
 0  Date             150000 non-null   object 
 1  Time             150000 non-null   object 
 2  Booking ID       150000 non-null   object 
 ...

```

### What this tells us:

- **RangeIndex:** Row numbers (0 to 149999)
- **150000 entries:** Total number of rows
- **21 columns:** Total number of columns
- **Non-Null Count:** How many non-missing values in each column
  - If less than 150000, that column has missing data
- **Dtype:** Data type of each column
  - object = Text/String
  - int64 = Whole numbers
  - float64 = Decimal numbers

- `datetime64` = Date/Time

## Key Things to Look For:

1. **Missing values:** Columns with Non-Null Count < 150000
  2. **Data types:** Are they correct? (dates should be datetime, numbers should be numeric)
  3. **Memory usage:** How much RAM the DataFrame uses
- 

## Step 2.3: Get Column Names



python

```
# List all column names
print("Column names in the dataset:")
print(df.columns.tolist())
```

## Understanding `df.columns`:

- **What it is:** An Index object containing all column names
- **.tolist():** Converts it to a regular Python list
- **Why useful:**
  - Check exact spelling of columns (no typos!)
  - Copy-paste column names for filtering
  - Understand data structure



python

```
# Better formatted view
print("\nColumns grouped by category:")
print("\n📅 Time Information:")
print("- Date, Time")
print("\n🚗 Booking Information:")
print("- Booking ID, Booking Status, Vehicle Type")
print("\n🔑 Location Information:")
print("- Pickup Location, Drop Location")
print("\n💰 Financial Information:")
print("- Booking Value, Payment Method")
print("\n⭐ Rating Information:")
print("- Driver Ratings, Customer Rating")
print("\n✖ Cancellation Information:")
print("- Cancelled Rides by Customer, Reason for cancelling by Customer")
print("- Cancelled Rides by Driver, Driver Cancellation Reason")
print("- Incomplete Rides, Incomplete Rides Reason")
```

---

## Step 2.4: Basic Statistics



python

```
# Get statistical summary of numerical columns
df.describe()
```

### Understanding df.describe() Output:



|       | Booking Value | Ride Distance | Driver Ratings | Customer Rating |
|-------|---------------|---------------|----------------|-----------------|
| count | 150000.000    | 150000.000    | 145000.000     | 148000.000      |
| mean  | 387.45        | 15.32         | 4.12           | 4.23            |
| std   | 125.67        | 8.45          | 0.89           | 0.75            |
| min   | 50.00         | 0.50          | 1.00           | 1.00            |
| 25%   | 295.00        | 9.00          | 3.50           | 4.00            |
| 50%   | 380.00        | 14.00         | 4.20           | 4.30            |
| 75%   | 475.00        | 20.00         | 4.80           | 4.70            |
| max   | 1250.00       | 95.00         | 5.00           | 5.00            |

## What Each Row Means:

- **count:** Number of non-null values (if less than total, some are missing)
- **mean:** Average value
- **std:** Standard deviation (how spread out the values are)
  - High std = values are very different from each other
  - Low std = values are similar
- **min:** Smallest value
- **25%:** First quartile - 25% of values are below this
- **50%:** Median (middle value) - 50% of values are below this
- **75%:** Third quartile - 75% of values are below this
- **max:** Largest value

## What to Look For:

1. **Unrealistic values:**
  - Negative prices?
  - Ratings above 5?
  - Distance of 0 km?
2. **Outliers:**
  - Max value way larger than 75%?
  - Min value way smaller than 25%?
3. **Missing data:**
  - Count less than total rows?



python

```
# Get statistics for specific columns
print("\nBooking Value Statistics:")
print(df['Booking Value'].describe())

print("\nDriver Ratings Statistics:")
print(df['Driver Ratings'].describe())
```

## Step 2.5: Understanding Data Types



python

```
# Check data types of all columns
print("Data types of each column:")
print(df.dtypes)
```

### Common Data Types in Pandas:

| Dtype      | Meaning                | Example                |
|------------|------------------------|------------------------|
| object     | Text/String            | "Delhi", "Cash", "SUV" |
| int64      | Integer (whole number) | 1, 42, 150000          |
| float64    | Decimal number         | 3.14, 387.45, 4.2      |
| bool       | True/False             | True, False            |
| datetime64 | Date and time          | 2022-01-15 14:30:00    |

### Why Data Types Matter:

- You can't do math on object type (even if it looks like a number)
- Dates stored as object can't be used for time-based analysis
- Wrong data types cause errors and prevent analysis



python

```
# Check specific column type
print(f'Date column type: {df['Date'].dtype}')
print(f'Booking Value column type: {df['Booking Value'].dtype}')
```

---

## Part 3: Initial Data Quality Check (15 minutes)

### Step 3.1: Check for Missing Values



python

```

# Count missing values in each column
print("Missing values per column:")
missing_values = df.isnull().sum()
print(missing_values)

# Show only columns with missing values
print("\nColumns with missing values:")
missing_cols = missing_values[missing_values > 0]
print(missing_cols)

# Calculate percentage of missing values
print("\nPercentage of missing values:")
missing_percent = (missing_values / len(df)) * 100
print(missing_percent[missing_percent > 0])

```

### Understanding `df.isnull()`:

- **What it does:** Returns a DataFrame of True/False values
  - True where value is missing (NaN, None, null)
  - False where value exists
- `.sum()`: Counts the True values (missing data) per column
- **Why chain them:** `df.isnull().sum()` gives count of missing values per column

### Example Output:



|                                   |       |
|-----------------------------------|-------|
| Date                              | 0     |
| Time                              | 0     |
| Booking ID                        | 0     |
| Cancelled Rides by Customer       | 85000 |
| Reason for cancelling by Customer | 85000 |
| Driver Ratings                    | 5000  |

### What This Tells Us:

- Some columns have NO missing values (Date, Time, Booking ID)
- Cancellation columns have many missing values (this is expected - not all rides are cancelled!)
- Driver Ratings has 5000 missing values (3.33% of data)

## Step 3.2: Check for Duplicate Rows



python

```
# Check for duplicate rows
duplicate_count = df.duplicated().sum()
print(f"Number of duplicate rows: {duplicate_count}")

# Show duplicate rows if any
if duplicate_count > 0:
    print("\nDuplicate rows:")
    duplicates = df[df.duplicated()]
    print(duplicates.head())
```

### Understanding `df.duplicated()`:

- **What it does:** Returns True for each duplicate row (after the first occurrence)
- `.sum()`: Counts how many duplicates exist
- **Example:**



Row 1: A, B, C <- Not marked as duplicate  
Row 2: X, Y, Z <- Not marked as duplicate  
Row 3: A, B, C <- Marked as duplicate (same as Row 1)

### Parameters you can use:



python

```
# Check duplicates based on specific columns
df.duplicated(subset=['Booking ID']) # Only check if Booking ID is duplicate

# Keep different occurrences
df.duplicated(keep='first') # Mark 2nd, 3rd, etc. as duplicates (default)
df.duplicated(keep='last') # Mark 1st, 2nd, etc. as duplicates
df.duplicated(keep=False) # Mark ALL duplicates (including first occurrence)
```

### Step 3.3: Unique Values Check



python

```
# Check unique values in categorical columns
categorical_columns = [
    'Booking Status',
    'Vehicle Type',
    'Payment Method'
]

print("Unique values in categorical columns:\n")
for col in categorical_columns:
    unique_count = df[col].nunique()
    unique_values = df[col].unique()
    print(f'{col}:')
    print(f' - Number of unique values: {unique_count}')
    print(f' - Values: {unique_values}')
    print()
```

#### Understanding .nunique() vs .unique():

- `.nunique()`: Returns COUNT of unique values (just a number)
- `.unique()`: Returns ARRAY of actual unique values

#### Example Output:



Booking Status:

- Number of unique values: 3
- Values: ['Success' 'Cancelled by Customer' 'Cancelled by Driver']

Vehicle Type:

- Number of unique values: 4
- Values: ['Auto' 'Bike' 'Prime Sedan' 'Prime SUV']

#### Why This Matters:

- Helps identify typos or inconsistent naming ("sedan" vs "Sedan" vs "SEDAN")
- Reveals all possible categories for analysis
- Helps spot data quality issues



python

```
# Value counts - how many times each unique value appears
print("Frequency of each Booking Status:")
print(df['Booking Status'].value_counts())

print("\nFrequency of each Vehicle Type:")
print(df['Vehicle Type'].value_counts())
```

### Understanding .value\_counts():

- **What it does:** Counts how many times each unique value appears
- **Returns:** A Series sorted by frequency (most common first)
- **Useful for:** Understanding distribution of categorical data

### Example Output:



```
Booking Status
Success           120000
Cancelled by Customer   20000
Cancelled by Driver     10000
Name: Booking Status, dtype: int64
```

---

## Part 4: Basic Analysis Questions (25 minutes)

Now let's answer some simple questions about our data!

### Question 1: How many total bookings do we have?



python

```
# Method 1: Using shape
total_bookings = df.shape[0]
print(f"Total bookings: {total_bookings:,}")
```

```
# Method 2: Using len()
total_bookings = len(df)
print(f"Total bookings: {total_bookings:,}")
```

```
# Method 3: Counting non-null Booking IDs
total_bookings = df['Booking ID'].count()
print(f"Total bookings: {total_bookings:,}")
```

## Which Method to Use?

- `df.shape[0]`: Fastest, always use this for row count
- `len(df)`: Also fast, more Pythonic
- `.count()`: Only use when you want to count non-null values

---

## Question 2: What's the total revenue generated?



python

```
# Sum all booking values
total_revenue = df['Booking Value'].sum()
print(f"Total revenue: ₹{total_revenue:.2f}")
```

```
# Average booking value
avg_booking = df['Booking Value'].mean()
print(f"Average booking value: ₹{avg_booking:.2f}")
```

```
# Median booking value
median_booking = df['Booking Value'].median()
print(f"Median booking value: ₹{median_booking:.2f}")
```

## Understanding `.sum()`, `.mean()`, `.median()`:

| Function               | What it calculates             | When to use                    |
|------------------------|--------------------------------|--------------------------------|
| <code>.sum()</code>    | Total (adds all values)        | Revenue, total distance, count |
| <code>.mean()</code>   | Average (sum / count)          | General "typical" value        |
| <code>.median()</code> | Middle value (50th percentile) | When data has outliers         |

## Mean vs Median Example:



Values: [10, 20, 30, 40, 1000]

Mean: 220 (skewed by the 1000)

Median: 30 (the middle value - more representative)

## Formatting Numbers:

- `:,:`: Adds thousand separators (58116750 → 58,116,750)
- `.2f`: Shows 2 decimal places (387.456789 → 387.46)
- `,.2f`: Both! (58116750.456 → 58,116,750.46)

---

## Question 3: How many unique customers do we have?



python

```
# Count unique customers
unique_customers = df['Customer ID'].nunique()
print(f'Number of unique customers: {unique_customers:,}')
```

```
# Average bookings per customer
avg_bookings_per_customer = len(df) / unique_customers
print(f'Average bookings per customer: {avg_bookings_per_customer:.2f}')
```

## Business Insight:

- If avg bookings per customer is high (>5), customers are loyal/frequent users
- If low (<2), mostly one-time users - need retention strategies

---

## Question 4: What are the most popular vehicle types?



python

```
# Count bookings by vehicle type
vehicle_counts = df['Vehicle Type'].value_counts()
print("Bookings by Vehicle Type:")
print(vehicle_counts)

# As percentages
vehicle_percentages = df['Vehicle Type'].value_counts(normalize=True) * 100
print("\nPercentage distribution:")
print(vehicle_percentages)
```

### Understanding .value\_counts() Parameters:

- **normalize=True**: Returns proportions (0.0 to 1.0) instead of counts
- **dropna=False**: Include missing values in count
- **sort=False**: Don't sort by frequency

### Example Output:



Vehicle Type

|             |       |       |
|-------------|-------|-------|
| Prime Sedan | 60000 | (40%) |
| Auto        | 45000 | (30%) |
| Prime SUV   | 30000 | (20%) |
| Bike        | 15000 | (10%) |

---

### Question 5: What's the completion rate?



python

```

# Count each booking status
status_counts = df['Booking Status'].value_counts()
print("Booking Status Distribution:")
print(status_counts)

# Calculate completion rate
total_rides = len(df)
completed_rides = status_counts['Success'] # or however success is labeled
completion_rate = (completed_rides / total_rides) * 100

print(f"\nCompletion Rate: {completion_rate:.2f}%")
print(f'Cancellation Rate: {100 - completion_rate:.2f}%')

```

### Business Metrics:

- **Completion Rate:** % of rides that finished successfully
- **Cancellation Rate:** % of rides cancelled (customer + driver)
- **Industry Standard:** Usually aim for >85% completion rate

## Part 5: Simple Visualizations (15 minutes)

### Visualization 1: Distribution of Booking Values



python

```

# Create histogram
plt.figure(figsize=(10, 6)) # Set size: width=10 inches, height=6 inches
plt.hist(df['Booking Value'], bins=50, edgecolor='black', color='skyblue')
plt.title('Distribution of Booking Values', fontsize=16, fontweight='bold')
plt.xlabel('Booking Value (₹)', fontsize=12)
plt.ylabel('Frequency (Number of Rides)', fontsize=12)
plt.grid(True, alpha=0.3) # Add grid with 30% opacity
plt.tight_layout() # Adjust spacing
plt.show()

```

### Understanding plt.hist() Parameters:

- **bins=50:** Number of bars in histogram (more bins = more detail)
- **edgecolor='black':** Border color around each bar
- **color='skyblue':** Fill color of bars
- **alpha=0.7:** Transparency (0=invisible, 1=solid)

## Understanding plt.figure() Parameters:

- **figsize=(10, 6)**: Size in inches (width, height)
- Standard sizes:
  - Small: (8, 6)
  - Medium: (10, 6)
  - Large: (12, 8)
  - Wide: (15, 6)

## What This Chart Shows:

- Most common booking values
- Whether prices are clustered or spread out
- If there are any unusual price points

## Visualization 2: Vehicle Type Distribution



python

```
# Create bar chart
vehicle_counts = df['Vehicle Type'].value_counts()

plt.figure(figsize=(10, 6))
vehicle_counts.plot(kind='bar', color='coral', edgecolor='black')
plt.title('Number of Rides by Vehicle Type', fontsize=16, fontweight='bold')
plt.xlabel('Vehicle Type', fontsize=12)
plt.ylabel('Number of Rides', fontsize=12)
plt.xticks(rotation=45, ha='right') # Rotate x-axis labels 45 degrees
plt.grid(True, alpha=0.3, axis='y') # Only horizontal grid lines
plt.tight_layout()
plt.show()
```

## Understanding plot() Parameters:

- **kind='bar'**: Type of plot
  - 'bar' - Vertical bars
  - 'barh' - Horizontal bars
  - 'line' - Line chart
  - 'hist' - Histogram
  - 'scatter' - Scatter plot
  - 'pie' - Pie chart

## Understanding plt.xticks() Parameters:

- **rotation=45**: Rotate labels 45 degrees (prevents overlap)
- **ha='right'**: Horizontal alignment (left, center, right)
- Common rotation values:

- 0 = Horizontal
  - 45 = Diagonal (most common)
  - 90 = Vertical
- 

## Visualization 3: Booking Status Pie Chart



python

```
# Create pie chart
status_counts = df['Booking Status'].value_counts()

plt.figure(figsize=(8, 8))
plt.pie(
    status_counts,
    labels=status_counts.index,
    autopct='%1.1f%%', # Show percentages with 1 decimal
    startangle=90,      # Start from top (90 degrees)
    colors=['#2ecc71', '#e74c3c', '#f39c12'], # Custom colors
    explode=(0.05, 0, 0) # "Explode" first slice slightly
)
plt.title('Distribution of Booking Status', fontsize=16, fontweight='bold')
plt.axis('equal') # Equal aspect ratio = circular pie
plt.tight_layout()
plt.show()
```

### Understanding plt.pie() Parameters:

- **autopct='%.1f%%'**: Format for percentages
    - %.1f = 1 decimal place (25.5%)
    - %1.0f = No decimals (25%)
    - %1.2f = 2 decimals (25.50%)
  - **startangle=90**: Where to start drawing (0-360 degrees)
  - **explode=(0.05, 0, 0)**: Pull slices out
    - One value per slice
    - 0 = Not pulled out
    - 0.1 = Pulled out 10%
  - **colors**: List of colors for each slice
    - Can use names: 'red', 'blue'
    - Or hex codes: '#2ecc71', '#e74c3c'
-

## Part 6: Wrap-up & Homework (5 minutes)

### Key Takeaways from Session 1:

#### What We Learned:

1. How to load data with `pd.read_csv()`
2. How to explore data with `.head()`, `.tail()`, `.info()`, `.describe()`
3. How to check data quality (missing values, duplicates)
4. How to get basic statistics (`.sum()`, `.mean()`, `.value_counts()`)
5. How to create simple visualizations

#### Key Functions Mastered:

- `df.head() / df.tail()` - View sample data
- `df.info()` - Get column information
- `df.describe()` - Get statistical summary
- `df.isnull().sum()` - Check missing values
- `df['column'].value_counts()` - Count unique values
- `df['column'].sum() / .mean() / .median()` - Basic math
- `plt.hist() / .plot() / .pie()` - Basic visualizations

### Homework for Next Session:



python

```

# 1. Find the most expensive ride
most_expensive = df['Booking Value'].max()
print(f'Most expensive ride: ₹{most_expensive}')

# 2. Find the longest distance traveled
longest_distance = df['Ride Distance'].max()
print(f'Longest distance: {longest_distance} km')

# 3. What's the most common payment method?
most_common_payment = df['Payment Method'].value_counts().index[0]
print(f'Most common payment method: {most_common_payment}')

# 4. How many rides were longer than 20 km?
long_rides = df[df['Ride Distance'] > 20].shape[0]
print(f'Rides longer than 20 km: {long_rides}')

# 5. What's the average rating given to drivers?
avg_driver_rating = df['Driver Ratings'].mean()
print(f'Average driver rating: {avg_driver_rating:.2f}')

```

### Challenge Questions:

1. Which pickup location had the most bookings?
2. What percentage of rides used each payment method?
3. Create a bar chart showing top 10 pickup locations by number of rides

## Session 2: Data Cleaning & Preparation

**Duration:** 90 minutes

**Goal:** Clean the data and prepare it for analysis

### Part 1: Understanding Data Quality Issues (10 minutes)

#### What Makes Data "Dirty"?



python

```

# Let's investigate our data quality systematically
print("==" * 50)
print("DATA QUALITY REPORT")
print("==" * 50)

# 1. Missing Values Summary
print("\n1. MISSING VALUES:")
missing_summary = df.isnull().sum()
print(missing_summary[missing_summary > 0])

# 2. Data Type Issues
print("\n2. DATA TYPES:")
print(df.dtypes)

# 3. Duplicate Check
print(f"\n3. DUPLICATES: {df.duplicated().sum()} duplicate rows found")

# 4. Inconsistent Values
print("\n4. UNIQUE VALUES IN KEY COLUMNS:")
for col in ['Booking Status', 'Vehicle Type', 'Payment Method']:
    print(f'{col}: {df[col].unique()}'")

```

## Common Data Quality Issues:

1. **Missing Values** - Blank cells, NaN, None
2. **Wrong Data Types** - Dates stored as text, numbers as strings
3. **Duplicates** - Same row appears multiple times
4. **Inconsistent Formatting** - "sedan" vs "Sedan", extra spaces
5. **Outliers** - Unrealistic values (negative prices, 1000 km rides in city)
6. **Invalid Values** - Ratings above 5, negative distances

## Part 2: Handling Missing Data (25 minutes)

### Step 2.1: Analyze Missing Data Patterns



python

```

# Create a visual summary of missing data
import matplotlib.pyplot as plt

# Calculate missing percentages
missing_percent = (df.isnull().sum() / len(df)) * 100
missing_data = missing_percent[missing_percent > 0].sort_values(ascending=False)

# Plot missing data
plt.figure(figsize=(12, 6))
missing_data.plot(kind='bar', color='salmon', edgecolor='black')
plt.title('Percentage of Missing Values by Column', fontsize=16, fontweight='bold')
plt.xlabel('Column Name', fontsize=12)
plt.ylabel('Percentage Missing (%)', fontsize=12)
plt.xticks(rotation=45, ha='right')
plt.grid(True, alpha=0.3, axis='y')
plt.tight_layout()
plt.show()

# Print detailed summary
for col in missing_data.index:
    count = df[col].isnull().sum()
    percent = missing_percent[col]
    print(f'{col}: {count} missing ({percent:.2f}%)')

```

## Step 2.2: Understanding WHY Data is Missing

### Three Types of Missing Data:

1. **Missing Completely at Random (MCAR)**
  - No pattern to missingness
  - Example: Device randomly failed to record data
  - Strategy: Usually safe to drop or impute
2. **Missing at Random (MAR)**
  - Missingness depends on other variables
  - Example: High-value rides less likely to report distance
  - Strategy: Impute based on related variables
3. **Missing Not at Random (MNAR)**
  - Missingness is related to the missing value itself
  - Example: Unhappy customers don't give ratings
  - Strategy: Careful! Might need special handling



python

```
# Let's investigate cancellation columns
print("Understanding Cancellation Data:")
print(f"Total rides: {len(df)}")
print(f"Missing 'Cancelled by Customer': {df['Cancelled Rides by Customer'].isnull().sum()}")
print(f"Missing 'Reason for cancelling': {df['Reason for cancelling by Customer'].isnull().sum()}")

# These nulls are INFORMATIVE - they mean "no cancellation happened"
# This is MNAR - the data is missing BECAUSE there was no cancellation
```

### Key Insight About Cancellation Data:

- Null values in cancellation columns mean "didn't cancel"
- These are NOT errors - they're informative!
- We should fill them with 0 or "Not Cancelled", not delete them

---

## Step 2.3: Strategy for Handling Missing Data



python

```
# Create a copy to work with (always preserve original!)
df_clean = df.copy()

print("Original shape:", df.shape)
```

### Why Create a Copy?

- `df.copy()` creates a completely separate DataFrame
- Changes to `df_clean` won't affect original `df`
- Lets you compare before/after
- Easy to restart if you make mistakes

### Without `.copy()`:



python

```
df_clean = df # This is just a reference, not a copy!
df_clean['column'] = 0 # This ALSO changes df!
```

## Step 2.4: Handling Cancellation Columns



python

```
# Strategy: Fill cancellation nulls with 0 (meaning "didn't happen")
cancellation_columns = [
    'Cancelled Rides by Customer',
    'Cancelled Rides by Driver',
    'Incomplete Rides'
]

for col in cancellation_columns:
    # Before
    before_nulls = df_clean[col].isnull().sum()

    # Fill nulls with 0
    df_clean[col] = df_clean[col].fillna(0)

    # After
    after_nulls = df_clean[col].isnull().sum()

    print(f'{col}:')
    print(f' Before: {before_nulls:,} nulls')
    print(f' After: {after_nulls:,} nulls')
    print(f' ✅ Filled {before_nulls:,} missing values with 0\n')
```

### Understanding .fillna():

- **What it does:** Replaces NaN/null values with specified value
- **Syntax:** `df['column'].fillna(value)`
- **Common fill values:**
  - 0 - For counts, flags
  - 'Unknown' - For categorical text
  - `.mean()` - Average of column
  - `.median()` - Middle value of column
  - `.mode()[0]` - Most common value
- **Returns:** New Series/DataFrame (doesn't change original unless `inplace=True`)



python

```
# Handle reason columns - fill with 'Not Applicable'  
reason_columns = [  
    'Reason for cancelling by Customer',  
    'Driver Cancellation Reason',  
    'Incomplete Rides Reason'  
]
```

```
for col in reason_columns:  
    before_nulls = df_clean[col].isnull().sum()  
    df_clean[col] = df_clean[col].fillna('Not Applicable')  
    after_nulls = df_clean[col].isnull().sum()  
  
    print(f'{col}:')  
    print(f' Before: {before_nulls:,} nulls')  
    print(f' After: {after_nulls:,} nulls')  
    print(f' ✅ Filled with "Not Applicable"\n')
```

---

## Step 2.5: Handling Rating Columns



python

```

# Check rating distributions before imputation
print("Driver Ratings Distribution:")
print(df_clean['Driver Ratings'].describe())

print("\nCustomer Rating Distribution:")
print(df_clean['Customer Rating'].describe())

# Strategy: Fill with MEDIAN (more robust than mean for ratings)
print("\nFilling missing ratings with median...")

# Driver Ratings
driver_median = df_clean['Driver Ratings'].median()
print(f"Driver Ratings median: {driver_median}")
df_clean['Driver Ratings'] = df_clean['Driver Ratings'].fillna(driver_median)

# Customer Rating
customer_median = df_clean['Customer Rating'].median()
print(f"Customer Rating median: {customer_median}")
df_clean['Customer Rating'] = df_clean['Customer Rating'].fillna(customer_median)

# Verify
print(f"\n✓ Driver Ratings nulls remaining: {df_clean['Driver Ratings'].isnull().sum()}")
print(f"✓ Customer Rating nulls remaining: {df_clean['Customer Rating'].isnull().sum()}")

```

## Mean vs Median for Imputation:

### Use MEAN when:

- Data is normally distributed (bell curve)
- No extreme outliers
- Example: Heights, test scores

### Use MEDIAN when:

- Data has outliers
- Skewed distribution
- Ratings (1-5 scale with clustering)

### Example:



Ratings: [4.5, 4.7, 4.8, 4.2, 4.6, 1.0, 4.3]

Mean: 4.0 (pulled down by the 1.0)

Median: 4.5 (more representative)

---

## Step 2.6: Final Missing Value Check



python

```
# Verify all missing values are handled
print("=" * 50)
print("FINAL MISSING VALUES CHECK")
print("=" * 50)

remaining_nulls = df_clean.isnull().sum()
print(remaining_nulls[remaining_nulls > 0])

if remaining_nulls.sum() == 0:
    print("\n✅ SUCCESS! All missing values handled!")
else:
    print(f"\n⚠️ Still have {remaining_nulls.sum()} missing values to address")
```

---

## Part 3: Data Type Conversions (20 minutes)

### Step 3.1: Converting Date Column



python

```

# Check current date type
print(f'Current Date type: {df_clean['Date'].dtype}')
print(f'Sample values:\n{df_clean['Date'].head()}')


# Convert to datetime
df_clean['Date'] = pd.to_datetime(df_clean['Date'])

# Verify conversion
print(f'\nNew Date type: {df_clean['Date'].dtype}')
print(f'Sample values:\n{df_clean['Date'].head()}')

```

## Understanding `pd.to_datetime()`:

- **What it does:** Converts strings to datetime objects
- **Why important:** Enables time-based operations
  - Extract day, month, year
  - Filter by date ranges
  - Calculate time differences
  - Sort chronologically

## Common Parameters:



python

```

# Specify format for faster parsing
pd.to_datetime(df['Date'], format='%Y-%m-%d')

# Handle errors
pd.to_datetime(df['Date'], errors='coerce') # Invalid dates become NaT
pd.to_datetime(df['Date'], errors='ignore') # Keep as string if fails

# Handle different date formats
pd.to_datetime(df['Date'], dayfirst=True) # DD/MM/YYYY
pd.to_datetime(df['Date'], yearfirst=True) # YYYY/MM/DD

```

## Date Format Codes:

- %Y - 4-digit year (2022)
- %y - 2-digit year (22)
- %m - Month as number (01-12)
- %d - Day as number (01-31)
- %H - Hour (00-23)
- %M - Minute (00-59)
- %S - Second (00-59)

## Step 3.2: Converting Time Column



python

```
# Check current time type
print(f'Current Time type: {df_clean["Time"].dtype}')
print(f"Sample values:\n{df_clean['Time'].head()}")

# Convert to datetime (even though it's just time)
df_clean['Time'] = pd.to_datetime(df_clean['Time'], format='%H:%M:%S').dt.time

# Verify
print(f"\nNew Time type: {df_clean['Time'].dtype}")
print(f"Sample values:\n{df_clean['Time'].head()}")
```

### Understanding .dt accessor:

- **What it is:** Special accessor for datetime operations
- **Only works on:** datetime64 columns
- **Allows you to extract:** date components, format dates, etc.

### Common .dt operations:



python

```
df['Date'].dt.year      # Extract year
df['Date'].dt.month     # Extract month (1-12)
df['Date'].dt.day       # Extract day (1-31)
df['Date'].dt.dayofweek # Day of week (0=Monday, 6=Sunday)
df['Date'].dt.day_name() # Day name ('Monday', 'Tuesday', etc.)
df['Date'].dt.month_name() # Month name ('January', 'February', etc.)
df['Date'].dt.quarter   # Quarter (1-4)
df['Date'].dt.weekofyear # Week number (1-52)
```

## Step 3.3: Extracting Date Components



python

```

# Extract useful date features for analysis
print("Extracting date components...")

# Day of week
df_clean['Day_of_Week'] = df_clean['Date'].dt.day_name()
print("✓ Created Day_of_Week column")

# Month
df_clean['Month'] = df_clean['Date'].dt.month_name()
print("✓ Created Month column")

# Day of month
df_clean['Day'] = df_clean['Date'].dt.day
print("✓ Created Day column")

# Hour from time
# Need to convert time back to datetime for extraction
df_clean['Hour'] = pd.to_datetime(df_clean['Time'].astype(str), format='%H:%M:%S').dt.hour
print("✓ Created Hour column")

# Is weekend?
df_clean['Is_Weekend'] = df_clean['Date'].dt.dayofweek >= 5 # Saturday=5, Sunday=6
print("✓ Created Is_Weekend column")

# Show sample
print("\nSample of new columns:")
print(df_clean[['Date', 'Time', 'Day_of_Week', 'Month', 'Hour', 'Is_Weekend']].head())

```

## Why Extract Date Components?

1. **Analysis by time periods:** "Which day has most bookings?"
2. **Seasonal patterns:** "Are Mondays busier than Fridays?"
3. **Time-based grouping:** "Revenue by month"
4. **Feature engineering:** Use in machine learning later

## Understanding Boolean Columns:



python

```
df_clean['Is_Weekend'] = df_clean['Date'].dt.dayofweek >= 5
```

- Returns True/False for each row
- dayofweek: 0=Monday, 1=Tuesday, ..., 5=Saturday, 6=Sunday
- >= 5 means Saturday or Sunday
- Can use in filtering: `df_clean[df_clean['Is_Weekend']]`

## Step 3.4: Creating Ride Outcome Column



python

```
# Create a categorical column for ride outcome
# This makes analysis easier than checking multiple columns

print("Creating Ride Outcome column...")

# Initialize with 'Completed'
df_clean['Ride_Outcome'] = 'Completed'

# Update based on cancellation columns
df_clean.loc[df_clean['Cancelled Rides by Customer'] > 0, 'Ride_Outcome'] = 'Cancelled by Customer'
df_clean.loc[df_clean['Cancelled Rides by Driver'] > 0, 'Ride_Outcome'] = 'Cancelled by Driver'
df_clean.loc[df_clean['Incomplete Rides'] > 0, 'Ride_Outcome'] = 'Incomplete'

# Verify
print("\nRide Outcome Distribution:")
print(df_clean['Ride_Outcome'].value_counts())
print(f"\n{checkmark} Created Ride_Outcome column")
```

### Understanding `.loc[]`:

- **What it does:** Accesses rows and columns by labels
- **Syntax:** `df.loc[row_condition, column_name]`
- **Use cases:**
  - Filter and update: `df.loc[condition, 'column'] = value`
  - Select subset: `df.loc[rows, columns]`
  - Boolean indexing: `df.loc[df['Age'] > 25]`

### Why Use `.loc` for Assignment?



python

```
# ❌ BAD - Can cause SettingWithCopyWarning
```

```
df[df['value'] > 100]['category'] = 'High'
```

```
# ✅ GOOD - Clear and explicit
```

```
df.loc[df['value'] > 100, 'category'] = 'High'
```

## Part 4: Handling Inconsistent Data (15 minutes)

### Step 4.1: Standardizing Text Columns



python

```
# Check for inconsistencies in categorical columns
```

```
print("Checking for inconsistencies...\n")
```

```
# Vehicle Type
```

```
print("Vehicle Type unique values:")
```

```
print(df_clean['Vehicle Type'].unique())
```

```
print(f'Count: {df_clean["Vehicle Type"].nunique()}')
```

```
# Check for common issues
```

```
print("\nChecking for:")
```

```
print("- Leading/trailing spaces")
```

```
print("- Inconsistent capitalization")
```

```
print("- Typos")
```

```
# Clean up spaces and standardize
```

```
df_clean['Vehicle Type'] = df_clean['Vehicle Type'].str.strip() # Remove leading/trailing spaces
```

```
df_clean['Vehicle Type'] = df_clean['Vehicle Type'].str.title() # Title Case
```

```
# Same for other categorical columns
```

```
text_columns = ['Pickup Location', 'Drop Location', 'Payment Method', 'Booking Status']
```

```
for col in text_columns:
```

```
    df_clean[col] = df_clean[col].str.strip()
```

```
    print(f"✅ Cleaned {col}")
```

**Understanding .str accessor:**

- **What it is:** Special accessor for string operations
- **Only works on:** Object (string) columns
- **Like:** Python string methods but for entire columns

### Common .str operations:



python

```
df['column'].str.lower()      # Convert to lowercase
df['column'].str.upper()      # Convert to UPPERCASE
df['column'].str.title()      # Convert To Title Case
df['column'].str.strip()       # Remove leading/trailing spaces
df['column'].str.replace('old', 'new') # Replace text
df['column'].str.contains('text') # Check if contains text
df['column'].str.len()         # Length of each string
df['column'].str.split(',')    # Split by delimiter
```

---

## Step 4.2: Fixing Data Entry Errors



python

```

# Example: Standardize location names
print("\nStandardizing location names...")

# Create mapping for common variations
location_mapping = {
    'Gurgaon': 'Gurugram', # Official name
    'Delhi Airport': 'Indira Gandhi International Airport',
    'CP': 'Connaught Place',
    'Noida Sec': 'Noida Sector'
}

# Apply mapping
for old_name, new_name in location_mapping.items():
    df_clean['Pickup Location'] = df_clean['Pickup Location'].str.replace(old_name, new_name)
    df_clean['Drop Location'] = df_clean['Drop Location'].str.replace(old_name, new_name)

print("✓ Location names standardized")

# Verify
print("\nTop Pickup Locations after cleaning:")
print(df_clean['Pickup Location'].value_counts().head(10))

```

## Understanding .replace():



```

# Single value replacement
df['column'].replace('old', 'new')

# Multiple replacements with dictionary
df['column'].replace({'old1': 'new1', 'old2': 'new2'})

# Regex patterns
df['column'].str.replace(r'\d+', '', regex=True) # Remove all numbers

```

---

# Part 5: Handling Outliers (15 minutes)

## Step 5.1: Detecting Outliers



python

```
# Statistical method: Values beyond 3 standard deviations
print("Detecting outliers in numerical columns..\n")

numerical_cols = ['Booking Value', 'Ride Distance']

for col in numerical_cols:
    # Calculate statistics
    mean = df_clean[col].mean()
    std = df_clean[col].std()

    # Define outlier thresholds
    lower_bound = mean - 3 * std
    upper_bound = mean + 3 * std

    # Find outliers
    outliers = df_clean[(df_clean[col] < lower_bound) | (df_clean[col] > upper_bound)]

    print(f'{col}:')
    print(f' Mean: {mean:.2f}')
    print(f' Std Dev: {std:.2f}')
    print(f' Lower Bound: {lower_bound:.2f}')
    print(f' Upper Bound: {upper_bound:.2f}')
    print(f' Number of outliers: {len(outliers),} ({len(outliers)/len(df_clean)*100:.2f}%)')
    print(f' Min outlier: {outliers[col].min():.2f}')
    print(f' Max outlier: {outliers[col].max():.2f}')
    print()
```

### Understanding Outlier Detection Methods:

#### 1. Standard Deviation Method (3-sigma rule):



python

```
lower = mean - 3*std  
upper = mean + 3*std  
outliers = data outside [lower, upper]
```

- Works well for normally distributed data
- 99.7% of data should be within 3 standard deviations

## 2. IQR Method (Interquartile Range):



python

```
Q1 = data.quantile(0.25)  
Q3 = data.quantile(0.75)  
IQR = Q3 - Q1  
lower = Q1 - 1.5*IQR  
upper = Q3 + 1.5*IQR
```

- More robust to extreme outliers
- Standard for box plots

## 3. Domain Knowledge:



python

```
# Business rules  
outliers = df[df['Booking Value'] < 0] # Negative prices impossible  
outliers = df[df['Rating'] > 5] # Ratings above 5 impossible
```

---

## Step 5.2: Visualizing Outliers



python

```

# Box plot to visualize outliers
fig, axes = plt.subplots(1, 2, figsize=(15, 6))

# Booking Value
axes[0].boxplot(df_clean['Booking Value'].dropna())
axes[0].set_title('Booking Value Distribution', fontsize=14, fontweight='bold')
axes[0].set_ylabel('Booking Value (₹)')
axes[0].grid(True, alpha=0.3)

# Ride Distance
axes[1].boxplot(df_clean['Ride Distance'].dropna())
axes[1].set_title('Ride Distance Distribution', fontsize=14, fontweight='bold')
axes[1].set_ylabel('Distance (km)')
axes[1].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

```

## Understanding Box Plots:



Maximum (upper whisker)

|

----+---- Q3 (75th percentile)

| |

| BOX | Median (50th percentile)

| |

----+---- Q1 (25th percentile)

|

Minimum (lower whisker)

○ ○ Outliers (dots beyond whiskers)

- **Box:** Contains middle 50% of data (Q1 to Q3)
- **Line in box:** Median
- **Whiskers:** Extend to  $1.5 \times \text{IQR}$  beyond box
- **Dots:** Outliers beyond whiskers

## Step 5.3: Handling Outliers



python

```
# Decision: Keep outliers but flag them for investigation
print("Flagging outliers for investigation...")

# Create outlier flags
df_clean['Is_High_Value'] = df_clean['Booking Value'] > df_clean['Booking Value'].quantile(0.95)
df_clean['Is_Long_Distance'] = df_clean['Ride Distance'] > df_clean['Ride Distance'].quantile(0.95)

# Check flagged rides
print(f"\nHigh value rides (top 5%): {df_clean['Is_High_Value'].sum()}")
print(f"Long distance rides (top 5%): {df_clean['Is_Long_Distance'].sum()}

# Investigate one outlier
print("\nSample high-value ride:")
sample_outlier = df_clean[df_clean['Is_High_Value']].iloc[0]
print(f"Booking Value: ₹{sample_outlier['Booking Value']} ")
print(f"Distance: {sample_outlier['Ride Distance']} km")
print(f"Vehicle Type: {sample_outlier['Vehicle Type']} ")
print(f"Pickup: {sample_outlier['Pickup Location']} ")
print(f"Drop: {sample_outlier['Drop Location']}")
```

### Strategies for Handling Outliers:

#### 1. Keep them (recommended if legitimate):



python

```
# Just flag for analysis
df['is_outlier'] = (df['value'] > threshold)
```

#### 2. Remove them (careful! losing data):



python

```
# Only if clearly errors  
df_clean = df[(df['value'] >= lower) & (df['value'] <= upper)]
```

### 3. Cap them (winsorization):



python

```
# Replace extreme values with threshold  
df['value'] = df['value'].clip(lower=lower, upper=upper)
```

### 4. Transform them (log, square root):



python

```
# Reduce impact of extremes  
df['value_log'] = np.log(df['value'] + 1)
```

---

## Part 6: Data Validation & Export (5 minutes)

### Step 6.1: Final Data Quality Check



python

```
print("==" * 60)
print("FINAL DATA QUALITY REPORT")
print("==" * 60)

# 1. Shape
print(f"\n1. Dataset Shape:")
print(f" Original: {df.shape}")
print(f" Cleaned: {df_clean.shape}")

# 2. Missing Values
print(f"\n2. Missing Values:")
total_nulls = df_clean.isnull().sum().sum()
if total_nulls == 0:
    print(" ✅ No missing values!")
else:
    print(f" ⚠️ {total_nulls} missing values remain")

# 3. Data Types
print(f"\n3. Data Types:")
print(f" Datetime columns: {df_clean.select_dtypes(include='datetime64').columns.tolist()}")
print(f" Numeric columns: {df_clean.select_dtypes(include=['int64', 'float64']).columns.tolist()}")
print(f" Text columns: {df_clean.select_dtypes(include='object').columns.tolist()}")

# 4. New Columns Created
new_columns = [col for col in df_clean.columns if col not in df.columns]
print(f"\n4. New Columns Created ({len(new_columns)}):")
for col in new_columns:
    print(f" - {col}")

# 5. Duplicates
print(f"\n5. Duplicate Rows:")
print(f" {df_clean.duplicated().sum()} duplicates")

print("\n" + "==" * 60)
print(" ✅ DATA CLEANING COMPLETE!")
print("==" * 60)
```

---

## Step 6.2: Save Cleaned Data



python

```
# Save cleaned dataset
output_file = 'data/ncr_ride_bookings_cleaned.csv'
df_clean.to_csv(output_file, index=False)

print(f"✅ Cleaned data saved to: {output_file}")
print(f" Rows: {df_clean.shape[0]}")
print(f" Columns: {df_clean.shape[1]}")
print(f" File size: {os.path.getsize(output_file) / (1024*1024):.2f} MB")
```

### Understanding `to_csv()` Parameters:

- `index=False`: Don't save row numbers as a column
- `sep=', '`: Use comma as delimiter (default)
- `encoding='utf-8'`: Character encoding
- `na_rep='NA'`: How to represent missing values
- `columns=['col1', 'col2']`: Save only specific columns
- `header=True`: Include column names (default)

---

## Part 7: Session 2 Wrap-up

### Key Takeaways:

#### ✅ Data Cleaning Skills Mastered:

1. Identifying and understanding missing data patterns
2. Filling missing values appropriately (`.fillna()`)
3. Converting data types (`.astype()`, `pd.to_datetime()`)
4. Extracting date components (`.dt` accessor)
5. Standardizing text data (`.str` methods)
6. Detecting and handling outliers
7. Creating new useful columns

#### ✅ Key Functions:

- `df.fillna()` - Fill missing values
- `pd.to_datetime()` - Convert to datetime
- `.dt.day_name()`, `.dt.month_name()`, `.dt.hour` - Extract date parts
- `.str.strip()`, `.str.title()`, `.str.replace()` - Clean text
- `.loc[]` - Conditional selection and assignment
- `.quantile()` - Find percentiles
- `df.to_csv()` - Save cleaned data

## Homework:



python

# 1. Check average booking value by vehicle type

```
avg_by_vehicle = df_clean.groupby('Vehicle Type')['Booking Value'].mean()  
print(avg_by_vehicle)
```

# 2. Find completion rate by day of week

```
completion_by_day = df_clean.groupby('Day_of_Week')['Ride_Outcome'].value_counts(normalize=True)  
print(completion_by_day)
```

# 3. Compare weekend vs weekday revenue

```
weekend_revenue = df_clean.groupby('Is_Weekend')['Booking Value'].sum()  
print(weekend_revenue)
```

# 4. Identify peak booking hours

```
peak_hours = df_clean['Hour'].value_counts().sort_index()  
print(peak_hours)
```

## Session 3: Exploratory Data Analysis (EDA)

**Duration:** 90 minutes

**Goal:** Discover patterns and insights in the data

### Part 1: Business Questions Framework (10 minutes)

#### The Questions We'll Answer:

##### 1. Revenue Analysis:

- Which locations generate most revenue?
- What's the average booking value by vehicle type?
- How does revenue vary by time of day?

##### 2. Cancellation Analysis:

- What's the cancellation rate?
- Why do customers cancel?
- Why do drivers cancel?
- Which locations have highest cancellations?

##### 3. Time-Based Patterns:

- Which days are busiest?
- What are peak booking hours?
- Weekend vs weekday patterns?

#### 4. Customer Satisfaction:

- Average ratings by vehicle type?
- Correlation between distance and ratings?
- Do longer rides get better ratings?

#### 5. Geographic Insights:

- Most popular pickup locations?
- Most common routes?
- Distance patterns by location?

---

## Part 2: Revenue Analysis (20 minutes)

### Question 1: Top Revenue-Generating Locations



python

```

# Group by location and calculate total revenue
revenue_by_location = df_clean.groupby('Pickup Location')['Booking Value'].agg([
    ('Total_Revenue', 'sum'),
    ('Avg_Booking', 'mean'),
    ('Num_Rides', 'count')
]).sort_values('Total_Revenue', ascending=False)

# Show top 10
print("Top 10 Revenue-Generating Locations:")
print(revenue_by_location.head(10))

# Visualize
plt.figure(figsize=(12, 6))
top_10_locations = revenue_by_location.head(10)
plt.barh(range(len(top_10_locations)), top_10_locations['Total_Revenue'], color='green')
plt.yticks(range(len(top_10_locations)), top_10_locations.index)
plt.xlabel('Total Revenue (₹)', fontsize=12)
plt.title('Top 10 Locations by Revenue', fontsize=16, fontweight='bold')
plt.gca().invert_yaxis() # Highest at top
plt.grid(True, alpha=0.3, axis='x')
plt.tight_layout()
plt.show()

```

## Understanding .agg() with Custom Names:



python

```
.agg([
    ('Custom_Name', 'function'),
    ('Another_Name', 'mean')
])
```

- Creates columns with specified names
- More readable than default names
- Can apply multiple functions to same column

## Understanding .gca().invert\_yaxis():

- `.gca()`: "Get Current Axes" - gets the current plot
- `.invert_yaxis()`: Flips y-axis (highest value at top)
- Useful for rankings - #1 should be at top!

